

LLMs for Social Scientists

A Progressive Journey from Fundamentals to Advanced Applications

Yunus Emre Tapan

2025-07-30

Table of contents

LLMs for Social Scientists	8
Preface	9
About This Book	9
Why This Book Exists	9
Who This Book Is For	9
How This Book Is Organized	9
Prerequisites	10
About the Author	10
About BLISS	10
Navigation	10
Getting Started	11
Getting Started	13
Course Overview	15
Learning Philosophy	15
Why This Matters for Social Science Research	15
Course Structure	15
Chapter 1: Fundamentals of LLMs	15
Chapter 2: LLMs in Action	15
Chapter 3: Contextualizing LLMs	15
Learning Approach	16
Progressive Complexity	16
Hands-On Practice	16
Research-Focused Examples	16
What You'll Need	16
Technical Prerequisites	16
Mindset	16
Getting the Most from This Course	16
Before Each Chapter	16
During Each Chapter	17
After Each Chapter	17
BLISS Workshop Connection	17
Navigation	17
Technical Setup & Learning Strategies	18
Overview	18

Technical Setup	18
Google Colab Environment	18
Getting Started	18
Essential Python Concepts	18
Learning Strategies for Social Scientists	18
Overcoming Math Anxiety	18
Effective Learning Techniques	19
Building Confidence	19
Common Challenges and Solutions	19
“I’m not good at math”	19
“I don’t have a computer science background”	19
“I don’t have time”	19
Getting Help When Stuck	19
Self-Help Strategies	19
Asking for Help	20
Connecting to Your Research	20
Throughout the Course	20
Research Applications to Explore	20
Navigation	20

Fundamentals of LLMs 21

Fundamentals of LLMs 22

Neural Networks Fundamentals 24

Learning Objectives	24
Introduction	24
What You Should Know	24
Prerequisites	24
Learning Approach	24
Session 1: Perceptron by Hand (50 minutes)	25
What we’ll do: Build your first “artificial brain cell”	25
The Problems You’ll Solve	25
Session 2: Backpropagation by Hand (50 minutes)	25
What we’ll do: Discover how networks learn from their mistakes	25
Neural Network Example**: Complex pattern recognition	25
Mathematical Skills You’ll Develop	25
Wrap-up and Reflection	26
Review the Complete Journey	26
Key Takeaways	26
Connection to Research	26
Why This Matters for Social Science	26
Research Applications	26
Navigation	26

Transformers and Embeddings Fundamentals 27

Learning Objectives	27
-------------------------------	----

Introduction	27
Pre-Chapter Learning (2 hours)	27
Materials:	27
Session 1: Vector Embeddings Deep Dive (30 minutes)	28
Understanding How Words Become Numbers	28
Vector Database Applications	28
Session 2: Self-Attention Mechanism by Hand (40 minutes)	28
The Heart of Transformer Architecture	28
Why Attention Matters	28
Session 3: Complete Transformer Architecture (25 minutes)	29
From Attention to Full Model	29
Transformer Components	29
Key Concepts	29
Connection to Research	29
Why Transformers Matter for Social Science	29
Research Applications	30
Advanced Topics Preview	30
Multi-Head Attention	30
Positional Encoding	30
Layer Normalization	30
Connection to Chapter 3	30
Navigation	30

LLMs in Action 31

LLMs in Action 32

LLM Capabilities and Limitations 34

LLM Capabilities and Limitations 35

Learning Objectives	35
Introduction	35
Pre-Chapter Learning (30 minutes)	35
Materials:	35
Session 1: Exploring LLM Capabilities (45 minutes)	35
What LLMs Can Do	35
Hands-On Exercise: Model Exploration	36
Session 2: Understanding LLM Limitations (45 minutes)	36
What LLMs Cannot Do	36
Exercise: Identifying Limitations	36
Session 3: Practical Applications for Research (30 minutes)	37
Research Use Cases	37
Exercise: Research Application	37
Ethical Considerations	37
Responsible Use Guidelines	37
Connection to Research	38
When to Use LLMs	38

Best Practices	38
Navigation	38
Working with LLM APIs	39
Working with LLM APIs	40
Learning Objectives	40
Introduction	40
Pre-Chapter Learning (30 minutes)	40
Materials:	40
Session 1: Understanding APIs (30 minutes)	40
What are APIs?	40
Key Concepts	41
Example API Call	41
Session 2: Prompt Engineering (45 minutes)	41
What is Prompt Engineering?	41
Basic Prompting Techniques	41
Advanced Prompting Techniques	42
Session 3: Practical Applications (45 minutes)	43
Research Applications	43
Exercise: Building a Research Tool	44
Error Handling and Best Practices	44
Common API Errors	44
Best Practices	44
Connection to Research	45
When to Use APIs vs. Local Models	45
Research Workflow Integration	45
Navigation	45
Contextualizing LLMs	46
Contextualizing LLMs	47
RAG and Context Engineering	49
RAG and Context Engineering	50
Learning Objectives	50
Introduction	50
Pre-Chapter Learning (30 minutes)	50
Materials:	50
Session 1: Understanding RAG (40 minutes)	50
What is Retrieval-Augmented Generation?	50
How RAG Works	51
Example: Philosophy Research	51
Advantages of RAG	51
Session 2: Vector Databases and Semantic Search (35 minutes)	51
Understanding Vector Embeddings	51
Semantic Search Process	51

Example: Research Paper Search	52
Vector Database Applications	52
Session 3: Knowledge Graphs (30 minutes)	52
What are Knowledge Graphs?	52
Knowledge Graph Structure	52
Example: Research Domain Knowledge Graph	52
Building Knowledge Graphs	53
Session 4: GraphRAG - Combining Approaches (25 minutes)	53
What is GraphRAG?	53
GraphRAG Process	53
Example: Philosophy Research with GraphRAG	53
Advantages of GraphRAG	54
Practical Applications for Research	54
Literature Review Enhancement	54
Survey Analysis	54
Research Planning	54
Implementation Considerations	54
Technical Requirements	54
Best Practices	55
Connection to Research	55
When to Use Each Approach	55
Research Workflow Integration	55
Navigation	55

LLM Showcase and Interface Design 56

LLM Showcase and Interface Design 57

Learning Objectives	57
Introduction	57
Pre-Chapter Learning (90 minutes)	57
Required Materials:	57
Session 1: Interface Design Principles (30 minutes)	57
Why Interface Design Matters for AI	57
Key Design Principles	58
Session 2: Building with Gradio (45 minutes)	58
What is Gradio?	58
Basic Gradio Interface	58
Advanced Gradio Features	59
Session 3: Research Application Examples (45 minutes)	59
Example 1: Survey Analysis Tool	59
Example 2: Literature Review Assistant	60
Example 3: Interview Analysis Tool	60
Session 4: Best Practices for AI Interfaces (30 minutes)	60
Transparency and Explainability	60
User Experience Considerations	60
Research-Specific Considerations	60
Implementation Guide	61
Step 1: Define Your Application	61
Step 2: Build the Backend	61

Step 3: Create the Interface	61
Step 4: Deploy and Share	61
Connection to Research	61
When to Build Interfaces	61
Best Practices for Research Interfaces	62
Navigation	62

References	63
-------------------	-----------

References	64
-------------------	-----------

LLMs for Social Scientists

A Progressive Journey from Fundamentals to Advanced Applications

Preface

About This Book

This book represents the collective knowledge from the **Boston LLMs Initiative for Social Scientists (BLISS)** workshop series. Our mission is to make Large Language Models accessible and useful for social science researchers, while fostering critical discussions about their limitations and biases.

Why This Book Exists

Large Language Models are transforming how we conduct research, analyze text, and interact with data. However, many social science researchers find these tools intimidating or inaccessible. This book bridges that gap by providing:

- **Progressive learning** from fundamentals to advanced applications
- **Social science focus** with research-relevant examples
- **Hands-on practice** with real-world applications
- **Critical perspective** on limitations and ethical considerations

Who This Book Is For

This book is designed for social science researchers who want to understand and apply LLMs to their research workflows. Whether you're a graduate student, faculty member, or research professional, this book provides the foundation you need to confidently use LLMs in your work.

How This Book Is Organized

The book follows a progressive structure:

Getting Started: Foundation and preparation **Fundamentals:** Understanding how LLMs work **LLMs in Action:** Practical applications and APIs **Contextualizing LLMs:** Advanced techniques and complete applications

Each chapter builds upon the previous one, ensuring you have a solid foundation before moving to more advanced topics.

Prerequisites

- Basic familiarity with Python (we'll provide refreshers)
- Willingness to work through mathematical concepts step-by-step
- Interest in applying AI to social science research

About the Author

Yunus Emre Tapan is the founder of BLISS and a researcher focused on making AI tools accessible to social scientists.

About BLISS

The **Boston LLMs Initiative for Social Scientists (BLISS)** organizes funded workshop series teaching social science researchers how to use large language models in their work, while discussing limitations and biases. Visit yemretapan.com/bliss for upcoming events and workshops.

This book is a living document that will be updated as the field evolves. We welcome feedback and contributions from the research community.

Navigation

Next: [Chapter 1: Getting Started](#) →

[Join the Discussion](#)

Have questions about this content? Want to share insights with other social science researchers? Join the conversation in our [GitHub Discussions](#).

[View Discussions](#) [Start New Discussion](#)

For Researchers: This discussion space is designed for social science researchers to share experiences, ask methodological questions, and discuss applications of LLMs in their work.

Getting Started

Report Content Issue Found an error or have a suggestion? Help improve this resource for fellow researchers.

Getting Started

This section provides the foundation for your journey into Large Language Models. You'll learn about the course structure, set up your technical environment, develop effective learning strategies, and understand how LLMs can be applied to social science research.

Course Overview

Learning Philosophy

This course is designed specifically for social science researchers who want to understand and use Large Language Models (LLMs) in their work. We take a progressive approach that builds from fundamental concepts to practical applications.

Why This Matters for Social Science Research

LLMs are transforming how we can analyze text, conduct research, and interact with data. Understanding these tools allows you to:

- **Analyze large text datasets** more efficiently
- **Generate research hypotheses** from literature reviews
- **Create survey instruments** and interview protocols
- **Process qualitative data** at scale
- **Communicate findings** more effectively

Course Structure

Chapter 1: Fundamentals of LLMs

Building the Foundation - Neural network basics and mathematical understanding - Transformer architecture and attention mechanisms - How modern LLMs process and generate text

Chapter 2: LLMs in Action

Practical Applications - Exploring LLM capabilities and limitations - Working with APIs and interfaces - Hands-on practice with real models

Chapter 3: Contextualizing LLMs

Advanced Applications - Retrieval-Augmented Generation (RAG) techniques - Knowledge graphs and advanced context methods - Building complete applications with user interfaces

Learning Approach

Progressive Complexity

Each chapter builds upon the previous one, ensuring you have a solid foundation before moving to more advanced topics.

Hands-On Practice

Theory is always paired with practical exercises, allowing you to immediately apply what you learn.

Research-Focused Examples

All examples and exercises are designed with social science research applications in mind.

What You'll Need

Technical Prerequisites

- Basic familiarity with Python (we'll provide refreshers)
- Access to Google Colab (free, no installation required)
- Willingness to work through mathematical concepts step-by-step

Mindset

- Curiosity about how AI tools work
- Patience with technical concepts
- Interest in applying new tools to research problems
- Openness to discussing limitations and biases

Getting the Most from This Course

Before Each Chapter

- Complete the pre-reading materials
- Set up your technical environment
- Review any prerequisites

During Each Chapter

- Work through exercises step-by-step
- Don't rush through mathematical concepts
- Ask questions when something isn't clear
- Reflect on how concepts apply to your research

After Each Chapter

- Reflect on what you learned
- Practice with your own examples
- Connect to the next chapter's content
- Consider attending BLISS workshops for hands-on practice

BLISS Workshop Connection

This book complements our in-person workshop series. While the book provides comprehensive self-paced learning, our workshops offer:

- **Hands-on guided practice** with expert instructors
- **Real-time feedback** and troubleshooting
- **Collaborative learning** with other researchers
- **Advanced topics** and specialized applications
- **Networking opportunities** with the BLISS community

Visit yemretapan.com/bliss for upcoming workshop dates and registration information.

Navigation

Next: [Technical Setup & Learning Strategies](#) →

Join the Discussion

Have questions about this content? Want to share insights with other social science researchers? Join the conversation in our [GitHub Discussions](#).

[View Discussions](#) [Start New Discussion](#)

For Researchers: This discussion space is designed for social science researchers to share experiences, ask methodological questions, and discuss applications of LLMs in their work.

Technical Setup & Learning Strategies

Overview

This chapter combines practical setup instructions with effective learning strategies. You'll learn how to get started with the technical tools and develop approaches that work for social science researchers.

Technical Setup

Google Colab Environment

- **No installation required** - everything runs in your browser
- **Free access** to powerful computing resources
- **Collaborative features** for sharing and version control

Getting Started

1. **Access:** Go to colab.research.google.com
2. **Sign in** with your Google account
3. **Create notebook** and start coding immediately

Essential Python Concepts

- **Variables:** Storing data (text, numbers, lists)
- **Functions:** Reusable code blocks
- **Libraries:** Pre-built tools for specific tasks
- **Data structures:** Lists, dictionaries for organizing information

Learning Strategies for Social Scientists

Overcoming Math Anxiety

- **Focus on concepts** over memorization
- **Real-world examples** make abstract ideas concrete
- **Step-by-step approach** - no rushing through complex topics
- **Practice with familiar data** from your research area

Effective Learning Techniques

- **Preview material** before diving in
- **Take notes** in your own words
- **Practice immediately** with exercises
- **Connect to your research** - think about applications
- **Use multiple resources** when needed

Building Confidence

- **Start small** and build gradually
- **Regular practice** (15-30 minutes daily)
- **Multiple learning styles** - visual, auditory, hands-on
- **Safe environment** - you can't break anything permanently

Common Challenges and Solutions

“I’m not good at math”

Solution: Focus on understanding concepts through examples. Many mathematical ideas are intuitive when explained properly.

“I don’t have a computer science background”

Solution: We explain technical concepts in social science terms. Many tools are designed to be user-friendly.

“I don’t have time”

Solution: Focus on concepts most relevant to your research. You don’t need to understand everything to use LLMs effectively.

Getting Help When Stuck

Self-Help Strategies

- **Re-read explanations** with fresh eyes
- **Look for concrete examples**
- **Break problems into smaller parts**
- **Try different approaches**

Asking for Help

- **Be specific** about what's confusing
- **Show your work** and thought process
- **Request examples** related to your research

Connecting to Your Research

Throughout the Course

- **Think about applications** to your specific research
- **Use your own examples** and data
- **Consider limitations** and ethical implications
- **Focus on practical skills** over theoretical perfection

Research Applications to Explore

- **Text analysis:** Interview transcripts, survey responses
- **Literature reviews:** Summarizing and synthesizing papers
- **Survey design:** Generating and analyzing questions
- **Data preprocessing:** Cleaning and preparing text data
- **Communication:** Writing summaries and reports

Navigation

Previous: [Course Overview](#) ←

Next: [Neural Networks Fundamentals](#) →

Fundamentals of LLMs

Fundamentals of LLMs

This section builds the theoretical foundation for understanding Large Language Models. You'll learn about neural networks, transformer architecture, and how modern LLMs process and generate text.

Neural Networks Fundamentals

Learning Objectives

By the end of this chapter, you will be able to: - Calculate perceptron outputs by hand for simple decisions - Understand neural network architecture and how layers connect - Perform forward propagation step-by-step through a network - Grasp backpropagation basics and how networks learn from mistakes - Visualize how networks process information - Prepare for advanced topics like transformers and embeddings

Introduction

Welcome to Neural Networks Fundamentals! Today you'll learn the mathematical foundations of deep learning by working through calculations by hand. No coding required - we'll focus on understanding how neural networks think and learn at the most basic level.

Duration: 2 hours

What you need: Basic algebra and willingness to work through math step-by-step

Approach: Hand calculations and visual understanding

What You Should Know

Prerequisites

- Basic algebra and arithmetic
- Understanding of functions (input \rightarrow output)
- Familiarity with graphs and coordinates
- Matrix multiplication basics (we'll review this)

Learning Approach

This lab is **calculation-intensive** but **beginner-friendly**: - **Step-by-step guidance**: Every calculation broken down clearly - **Real examples**: "Should we go to the beach?" and other relatable decisions - **Visual aids**: Diagrams showing information flow - **Immediate understanding**: See exactly how each step contributes

Session 1: Perceptron by Hand (50 minutes)

What we'll do: Build your first “artificial brain cell”

Exercise: Work through the “Should we go to the beach?” decision problem - Calculate weighted inputs step-by-step - Apply activation functions to make binary decisions - Understand the geometry of linear decision boundaries - Practice with multiple examples and variations - Explore different scenarios and weight adjustments

You'll learn: How the simplest neural network makes decisions

The Problems You'll Solve

Perceptron Example: “Should we go to the beach?” - **Inputs:** Weather conditions (sunny, temperature, humidity) - **Process:** Weight each factor and make a decision - **Output:** Yes/No decision with confidence

Session 2: Backpropagation by Hand (50 minutes)

What we'll do: Discover how networks learn from their mistakes

- Start with a simple multi-layer network example
- Make a prediction and calculate the error
- Work backwards through the network using chain rule
- Calculate gradients and weight updates step-by-step
- Understand gradient descent intuitively
- Practice with a complete learning cycle example

You'll learn: The magic behind how neural networks improve over time

Neural Network Example**: Complex pattern recognition

- **Architecture:** Input \rightarrow Hidden Layer \rightarrow Output
- **Forward Pass:** Calculate predictions
- **Backward Pass:** Learn from mistakes

Mathematical Skills You'll Develop

Mathematical Skills: - Matrix multiplication in neural network context - Understanding weighted sums and activation functions - Basic calculus concepts (derivatives for learning) - Visualization of mathematical concepts

Conceptual Skills: - How information flows through networks - Why networks can learn complex patterns - The relationship between structure and function - Preparing for more advanced architectures

Wrap-up and Reflection

Review the Complete Journey

- From simple decisions to learning
- Discuss what surprised you most about the mathematics
- Connect to modern AI: How these principles scale to large networks
- Preview Chapter 2: Transformers and Embeddings

Key Takeaways

- Neural networks are mathematical functions that can learn
- The perceptron is the building block of all neural networks
- Backpropagation enables learning through error correction
- These principles scale to modern AI systems

Connection to Research

Why This Matters for Social Science

- **Understanding AI tools:** Know how the tools you use actually work
- **Interpreting results:** Better understand what AI outputs mean
- **Designing studies:** Create better prompts and inputs for AI systems
- **Ethical considerations:** Understand limitations and potential biases

Research Applications

- **Survey design:** Understanding how AI might process responses
- **Data analysis:** Knowing how AI models interpret your data
- **Literature review:** Understanding AI summarization capabilities
- **Communication:** Explaining AI concepts to research participants

Navigation

Previous: [Technical Setup & Learning Strategies](#) ←

Next: [Transformers and Embeddings](#) →

Transformers and Embeddings Fundamentals

Learning Objectives

By the end of this chapter, you will be able to: - **Embedding Mastery:** Understand how words and concepts are represented as vectors - **Attention Mechanism:** Calculate self-attention by hand and understand its purpose - **Transformer Architecture:** Comprehend the complete transformer model structure - **Query-Key-Value System:** Master the fundamental attention computation framework

Introduction

This chapter explores the evolution from neural networks to transformers, the architecture that powers modern Large Language Models. You'll learn how words become numbers and how attention mechanisms enable models to understand context.

Pre-Chapter Learning (2 hours)

Materials:

1. **Vector Embeddings Foundation** (30 minutes)
 - [Word Embeddings Explained \(StatQuest\)](#) - 11-minute video on word representations
 - [What are Word Embeddings? \(IBM\)](#) - Technical overview (15 min read)
2. **Transformer Architecture** (60 minutes)
 - [Attention Is All You Need \(Paper Summary\)](#) - 30-minute explanation of the original paper
 - [Deep Dive into Transformers by Hand](#) - Detailed mathematical walkthrough (30 min read)
3. **Attention Mechanisms** (30 minutes)
 - [Keys, Queries, Values Explained](#) - Core attention concepts (15 min read)

Session 1: Vector Embeddings Deep Dive (30 minutes)

Understanding How Words Become Numbers

Key Concept: Vector embeddings transform words into numerical representations that capture meaning and relationships.

Exercise: Vector Embeddings Practice - Explore how words are represented as vectors - Understand similarity and distance in vector space - Practice with word analogies and relationships

You'll learn: How words become numbers that preserve meaning

Vector Database Applications

- **Similarity search:** Finding related concepts
- **Clustering:** Grouping similar ideas
- **Recommendation systems:** Suggesting related content
- **Research applications:** Literature review and synthesis

Session 2: Self-Attention Mechanism by Hand (40 minutes)

The Heart of Transformer Architecture

Key Concept: Self-attention allows positions to attend to other positions in the sequence, capturing relationships between words.

Exercise: Self-Attention Calculation - Calculate attention scores step-by-step - Understand query, key, and value matrices - Practice attention computation by hand - Visualize attention patterns

Mathematical Foundation:

$$\text{Attention}(Q,K,V) = \text{softmax}(QK^T/\sqrt{d_k})V$$

Why Attention Matters

- **Context understanding:** Words can attend to relevant context
- **Long-range dependencies:** Capture relationships across long sequences
- **Interpretability:** Attention weights show what the model focuses on
- **Flexibility:** Can attend to any position in the sequence

Session 3: Complete Transformer Architecture (25 minutes)

From Attention to Full Model

Exercise: Transformer Architecture Walkthrough - Understand the complete transformer structure - Explore encoder and decoder components - Practice with a complete example - Connect to real-world applications

Transformer Components

- **Input Embeddings:** Convert words to vectors
- **Positional Encoding:** Add position information
- **Multi-Head Attention:** Multiple attention mechanisms
- **Feed-Forward Networks:** Process attended information
- **Layer Normalization:** Stabilize training
- **Residual Connections:** Help with gradient flow

Key Concepts

Concept	Description	Mathematical Foundation
Vector Embeddings	Dense numerical representations of words/concepts	Continuous vector space mapping
Self-Attention	Mechanism allowing positions to attend to other positions	Scaled dot-product: $\text{Attention}(Q,K,V) = \frac{\exp(QK^T/\sqrt{d_k})}{\sum_j \exp(QK_j^T/\sqrt{d_k})}V$
Query, Key, Value	Three matrices that enable attention computation	Linear projections: $Q=XW_Q$, $K=XW_K$, $V=XW_V$

Connection to Research

Why Transformers Matter for Social Science

- **Text understanding:** Better comprehension of complex documents
- **Context awareness:** Models understand relationships between concepts
- **Scalability:** Can process large amounts of text efficiently
- **Interpretability:** Attention weights show what models focus on

Research Applications

- **Document analysis:** Understanding complex research papers
- **Survey analysis:** Processing open-ended responses
- **Literature review:** Synthesizing multiple sources
- **Content generation:** Creating research summaries

Advanced Topics Preview

Multi-Head Attention

- Multiple attention mechanisms running in parallel
- Captures different types of relationships
- Enables richer representations

Positional Encoding

- Adds position information to embeddings
- Enables understanding of word order
- Critical for sequence processing

Layer Normalization

- Stabilizes training process
- Improves convergence
- Essential for deep networks

Connection to Chapter 3

This chapter prepares you for Chapter 3 where you'll use pre-trained transformer models (BERT, GPT, etc.) through Hugging Face for practical NLP applications. You'll apply the theoretical understanding gained here to real-world tasks like text classification, NER, and generation.

Navigation

Previous: [Neural Networks Fundamentals](#) ←

Next: [LLM Capabilities](#) →

LLMs in Action

LLMs in Action

This section focuses on practical applications of Large Language Models. You'll explore LLM capabilities and limitations, work with APIs, and learn how to apply these tools to real research tasks.

LLM Capabilities and Limitations

LLM Capabilities and Limitations

Learning Objectives

By the end of this chapter, you will be able to: - Understand the capabilities and limitations of modern LLMs - Use Hugging Face to explore different models - Recognize when LLMs are appropriate for research tasks - Identify potential biases and limitations in LLM outputs - Apply LLMs to basic text analysis tasks

Introduction

This chapter introduces you to real Large Language Models through hands-on exploration. You'll learn what LLMs can and cannot do, and how to use them responsibly in your research.

Pre-Chapter Learning (30 minutes)

Materials:

1. **Hugging Face Introduction** (15 minutes)
 - [Hugging Face Course](#) - Introduction to the platform
 - [Model Hub Overview](#) - Understanding available models
2. **LLM Capabilities Overview** (15 minutes)
 - [What Can Language Models Do?](#) - Research paper on capabilities
 - [Limitations of Language Models](#) - Understanding limitations

Session 1: Exploring LLM Capabilities (45 minutes)

What LLMs Can Do

Text Generation - Complete sentences and paragraphs - Generate creative content - Answer questions based on training data - Translate between languages

Text Classification - Categorize documents by topic - Identify sentiment in text - Detect spam or inappropriate content - Classify research papers by field

Question Answering - Answer factual questions - Provide explanations - Generate summaries - Extract information from text

Code Generation - Write simple code snippets - Debug existing code - Explain code functionality - Generate documentation

Hands-On Exercise: Model Exploration

```
# Example: Exploring different models
from transformers import pipeline

# Text classification
classifier = pipeline("text-classification")
result = classifier("I love this research paper!")
print(result)

# Question answering
qa = pipeline("question-answering")
context = "The study found that social media use is associated with increased anxiety."
question = "What did the study find?"
answer = qa(question=question, context=context)
print(answer)
```

Session 2: Understanding LLM Limitations (45 minutes)

What LLMs Cannot Do

Factual Accuracy - May generate incorrect information - Cannot access real-time data - Limited to training data cutoff - May “hallucinate” plausible but false information

Reasoning and Logic - Struggle with complex reasoning - May make logical errors - Limited mathematical capabilities - Cannot perform true understanding

Bias and Fairness - Reflect biases in training data - May perpetuate stereotypes - Can amplify existing inequalities - Requires careful evaluation

Context Understanding - Limited context window - May miss nuanced meaning - Can be sensitive to prompt wording - May not understand domain-specific knowledge

Exercise: Identifying Limitations

Task: Analyze LLM outputs for potential issues - Test factual accuracy with known information - Evaluate bias in responses - Check for logical consistency - Assess domain knowledge limitations

Session 3: Practical Applications for Research (30 minutes)

Research Use Cases

Literature Review - Summarize research papers - Identify key themes and findings - Generate research questions - Create annotated bibliographies

Survey Analysis - Code open-ended responses - Identify common themes - Generate follow-up questions - Analyze sentiment in responses

Content Generation - Draft research summaries - Create presentation outlines - Generate grant proposal sections - Write accessible explanations

Data Preprocessing - Clean and standardize text data - Remove irrelevant content - Identify and handle missing data - Prepare data for analysis

Exercise: Research Application

```
# Example: Analyzing survey responses
responses = [
    "I feel overwhelmed by social media pressure",
    "Social media helps me stay connected",
    "I spend too much time on my phone"
]

# Use LLM to identify themes
for response in responses:
    # Analyze sentiment and themes
    # Generate coding categories
    # Identify key concepts
    pass
```

Ethical Considerations

Responsible Use Guidelines

Transparency - Document all LLM usage in methods - Disclose when AI tools are used - Be clear about limitations - Share prompts and parameters

Evaluation - Always verify LLM outputs - Cross-check with reliable sources - Consider multiple perspectives - Acknowledge potential biases

Privacy - Don't share sensitive data with LLMs - Be careful with participant information - Consider data retention policies - Protect confidential information

Fairness - Evaluate outputs for bias - Consider diverse perspectives - Test with different populations - Acknowledge limitations

Connection to Research

When to Use LLMs

Appropriate Uses - Initial exploration of large datasets - Generating research hypotheses - Creating draft content - Automating routine tasks

Inappropriate Uses - Making final research decisions - Replacing human judgment - Analyzing sensitive data - Generating definitive conclusions

Best Practices

For Literature Reviews - Use LLMs to identify themes and patterns - Always verify key findings manually - Cross-reference with original sources - Maintain critical perspective

For Survey Analysis - Use LLMs for initial coding - Verify coding with human review - Consider multiple coding approaches - Document coding decisions

For Content Generation - Use LLMs for drafting and brainstorming - Always edit and review generated content - Maintain your voice and perspective - Ensure accuracy and appropriateness

Navigation

Previous: [Transformers and Embeddings](#) ←

Next: [LLM APIs](#) →

Working with LLM APIs

Working with LLM APIs

Learning Objectives

By the end of this chapter, you will be able to: - Understand how to interact with LLM APIs - Use prompt engineering techniques effectively - Handle API responses and errors - Apply LLMs to real research tasks - Implement best practices for API usage

Introduction

This chapter teaches you how to work with Large Language Model APIs, including prompt engineering, response handling, and practical applications for research. You'll learn to communicate effectively with AI systems and get the most out of their capabilities.

Pre-Chapter Learning (30 minutes)

Materials:

1. **API Basics** (15 minutes)
 - [What are APIs?](#) - Introduction to APIs
 - [REST API Tutorial](#) - Understanding API communication
2. **Prompt Engineering** (15 minutes)
 - [Prompt Engineering Guide](#) - Best practices for prompts
 - [OpenAI API Documentation](#) - Technical reference

Session 1: Understanding APIs (30 minutes)

What are APIs?

Application Programming Interfaces (APIs) are ways for different software systems to communicate with each other. In the context of LLMs, APIs allow you to send text to an AI model and receive responses.

Key Concepts

HTTP Requests - **GET**: Retrieve information - **POST**: Send data to the server - **PUT/PATCH**: Update existing data - **DELETE**: Remove data

API Endpoints - Specific URLs that accept requests - Different endpoints for different functions - Authentication required for most services

Response Formats - JSON (JavaScript Object Notation) is most common - Structured data that's easy to parse - Contains the AI's response and metadata

Example API Call

```
import requests

# Example API call structure
url = "https://api.openai.com/v1/chat/completions"
headers = {
    "Authorization": "Bearer YOUR_API_KEY",
    "Content-Type": "application/json"
}
data = {
    "model": "gpt-3.5-turbo",
    "messages": [
        {"role": "user", "content": "Hello, how are you?"}
    ]
}

response = requests.post(url, headers=headers, json=data)
result = response.json()
print(result['choices'][0]['message']['content'])
```

Session 2: Prompt Engineering (45 minutes)

What is Prompt Engineering?

Prompt engineering is the practice of designing inputs to AI systems to get the best possible outputs. It's like learning to speak the AI's language effectively.

Basic Prompting Techniques

Clear Instructions

Bad: "Tell me about social media"

Good: "Provide a 3-paragraph summary of the psychological effects of social media use on teenagers"

Context Setting

Bad: "Analyze this text"

Good: "You are a social science researcher analyzing survey responses. Analyze the following"

Output Formatting

Bad: "List the themes"

Good: "Identify the top 3 themes in the following text and format your response as:

1. Theme name: Brief description
2. Theme name: Brief description
3. Theme name: Brief description"

Advanced Prompting Techniques

Few-Shot Learning

Example 1:

Input: "I feel stressed about work"

Output: Theme: Work-related stress

Example 2:

Input: "My relationship is causing anxiety"

Output: Theme: Relationship stress

Now analyze: "I'm worried about my finances"

Chain-of-Thought

Bad: "What's the answer?"

Good: "Let's approach this step by step:

1. First, identify the key concepts
2. Then, analyze their relationships
3. Finally, draw a conclusion

Question: [your question]"

Role-Based Prompting

"You are an expert social science researcher with 20 years of experience in qualitative data a

Session 3: Practical Applications (45 minutes)

Research Applications

Literature Review Assistance

```
prompt = """
You are a research assistant helping with a literature review on social media and mental health.

Given the following research paper abstract, please:
1. Identify the main research question
2. List the key findings
3. Note any limitations mentioned
4. Suggest how this relates to other research in the field

Abstract: [paper abstract]
"""
```

Survey Response Analysis

```
prompt = """
You are analyzing open-ended survey responses about social media use.

For each response, identify:
1. Primary theme (e.g., addiction, connection, privacy)
2. Sentiment (positive, negative, neutral, mixed)
3. Key concerns or benefits mentioned
4. Suggested follow-up questions

Response: [survey response]
"""
```

Content Generation

```
prompt = """
You are writing a research summary for a general audience.

Please convert the following academic findings into accessible language:
- Use simple, clear language
- Avoid jargon
- Include practical implications
- Keep it under 200 words

Findings: [research findings]
"""
```

Exercise: Building a Research Tool

Task: Create a simple survey analysis tool

```
def analyze_survey_response(response, api_client):  
    prompt = f"""  
    Analyze this survey response about social media use:  
  
    Response: {response}  
  
    Please provide:  
    1. Main theme (one word)  
    2. Sentiment (positive/negative/neutral)  
    3. Key concerns (list)  
    4. Suggested follow-up question  
    """  
  
    # Send to API and process response  
    # Return structured analysis  
    pass
```

Error Handling and Best Practices

Common API Errors

Rate Limiting - Too many requests too quickly - Solution: Implement delays between requests
- Monitor usage limits

Authentication Errors - Invalid or expired API key - Solution: Check API key and permissions - Rotate keys regularly

Model Errors - Model unavailable or overloaded - Solution: Implement retry logic - Have fallback models

Best Practices

Security - Never share API keys in code - Use environment variables - Monitor API usage and costs - Implement proper error handling

Efficiency - Batch requests when possible - Cache responses when appropriate - Use appropriate model sizes - Monitor response times

Quality - Always validate responses - Implement human review for important outputs - Test with diverse inputs - Document prompt strategies

Connection to Research

When to Use APIs vs. Local Models

Use APIs When: - You need the latest models - You don't have computational resources - You need quick prototyping - You want to try multiple models

Use Local Models When: - You have sensitive data - You need consistent performance - You want to avoid API costs - You need offline capabilities

Research Workflow Integration

Planning Phase - Use LLMs to generate research questions - Explore potential methodologies - Identify relevant literature

Data Collection - Generate survey questions - Create interview protocols - Design coding schemes

Analysis Phase - Code qualitative data - Identify themes and patterns - Generate hypotheses

Writing Phase - Draft research summaries - Create presentation outlines - Generate accessible explanations

Navigation

Previous: [LLM Capabilities](#) ←

Next: [RAG and Context](#) →

Contextualizing LLMs

Contextualizing LLMs

This section covers advanced techniques for providing context to Large Language Models. You'll learn about RAG, knowledge graphs, and how to build complete applications with user interfaces.

RAG and Context Engineering

RAG and Context Engineering

Learning Objectives

By the end of this chapter, you will be able to: - Understand the difference between Plain LLM, RAG, and GraphRAG approaches - Learn how vector databases enable semantic search - Build and interpret simple knowledge graphs - Practice model selection and fallback for robust AI research - Compare and reflect on the strengths of each approach for research

Introduction

This chapter explores advanced techniques for providing context to Large Language Models. You'll learn about Retrieval-Augmented Generation (RAG), knowledge graphs, and how to combine them for more accurate and contextual AI responses.

Pre-Chapter Learning (30 minutes)

Materials:

1. **Vector Databases Foundation** (15 minutes)
 - [What is a Vector Database? \(Pinecone\)](#) - Industry overview
 - [What is a Knowledge Graph? \(YouTube\)](#) - Visual intro
2. **RAG Concepts** (15 minutes)
 - [RAG vs. GraphRAG \(YouTube\)](#) - Quick comparison
 - Preview the hands-on exercises

Session 1: Understanding RAG (40 minutes)

What is Retrieval-Augmented Generation?

RAG combines the power of Large Language Models with external knowledge sources. Instead of relying solely on the model's training data, RAG retrieves relevant information and provides it as context to the LLM.

How RAG Works

1. **Query Processing:** Convert user question into search terms
2. **Retrieval:** Search external knowledge base for relevant documents
3. **Context Injection:** Provide retrieved documents to LLM
4. **Generation:** LLM generates answer using provided context

Example: Philosophy Research

```
# Example: Researching philosophical concepts
query = "What did Kant say about moral duty?"

# RAG Process:
# 1. Search philosophy papers for Kant + moral duty
# 2. Retrieve relevant passages
# 3. Provide context to LLM
# 4. Generate answer based on retrieved information
```

Advantages of RAG

Accuracy - Based on specific, relevant information - Reduces hallucination - More up-to-date than training data

Transparency - Can cite sources - Traceable to original documents - Verifiable information

Flexibility - Can use any knowledge base - Adaptable to different domains - Easy to update with new information

Session 2: Vector Databases and Semantic Search (35 minutes)

Understanding Vector Embeddings

Vector embeddings represent text as numerical vectors that capture meaning and relationships.

Semantic Search Process

1. **Document Embedding:** Convert documents to vectors
2. **Query Embedding:** Convert search query to vector
3. **Similarity Calculation:** Find most similar document vectors
4. **Retrieval:** Return most relevant documents

Example: Research Paper Search

```
# Example: Finding relevant research papers
query = "social media effects on mental health"

# Vector search finds papers about:
# - Social media and depression
# - Instagram and anxiety
# - Facebook and well-being
# - Digital technology and psychology
```

Vector Database Applications

Literature Review - Find related research papers - Identify research gaps - Discover emerging themes

Survey Analysis - Group similar responses - Identify common themes - Find representative examples

Content Recommendation - Suggest related articles - Find similar research questions - Recommend follow-up studies

Session 3: Knowledge Graphs (30 minutes)

What are Knowledge Graphs?

Knowledge graphs represent information as networks of connected entities and relationships. They capture not just individual facts, but the relationships between them.

Knowledge Graph Structure

Entities: People, places, concepts, events **Relationships:** Connections between entities

Properties: Attributes of entities

Example: Research Domain Knowledge Graph

Entities:

- Social Media (concept)
- Mental Health (concept)
- Instagram (platform)
- Depression (condition)

Relationships:

- Instagram -> affects -> Mental Health
- Social Media -> includes -> Instagram

- Depression -> is_a -> Mental Health condition
- Instagram -> associated_with -> Depression

Building Knowledge Graphs

Manual Construction - Define entities and relationships - Create structured data - Validate connections

Automated Extraction - Use NLP to extract entities - Identify relationships from text - Validate with human review

Hybrid Approaches - Start with manual structure - Automate extraction for scale - Human validation for quality

Session 4: GraphRAG - Combining Approaches (25 minutes)

What is GraphRAG?

GraphRAG combines traditional RAG with knowledge graph information to provide richer context and better understanding of relationships.

GraphRAG Process

1. **Traditional RAG:** Retrieve relevant documents
2. **Graph Enhancement:** Add knowledge graph information
3. **Relationship Analysis:** Identify connections between concepts
4. **Enhanced Context:** Provide both documents and graph information to LLM

Example: Philosophy Research with GraphRAG

```
# Research question: "How do different philosophers view free will?"

# Traditional RAG finds:
# - Papers about Kant and free will
# - Papers about Hume and determinism
# - Papers about contemporary views

# GraphRAG also finds:
# - Connections between Kant and Hume
# - Historical development of ideas
# - Related concepts (causality, responsibility)
# - Influence relationships between philosophers
```

Advantages of GraphRAG

Richer Context - Documents + relationships - Historical development - Influence networks - Related concepts

Better Understanding - Contextual relationships - Temporal development - Influence patterns - Conceptual connections

More Comprehensive Answers - Multiple perspectives - Historical context - Related concepts - Influence networks

Practical Applications for Research

Literature Review Enhancement

Traditional Approach - Read papers individually - Manual synthesis - Limited to direct citations

GraphRAG Approach - Automated paper discovery - Relationship identification - Comprehensive synthesis - Influence network analysis

Survey Analysis

Traditional Approach - Manual coding - Limited to direct responses - Time-consuming analysis

GraphRAG Approach - Automated theme identification - Relationship discovery - Comprehensive analysis - Pattern recognition

Research Planning

Traditional Approach - Manual literature review - Limited scope - Time-intensive

GraphRAG Approach - Automated gap identification - Comprehensive coverage - Relationship discovery - Efficient planning

Implementation Considerations

Technical Requirements

Vector Database - Pinecone, Weaviate, or similar - Storage for document embeddings - Search capabilities

Knowledge Graph - Neo4j, ArangoDB, or similar - Graph query capabilities - Relationship management

Integration - API connections - Data synchronization - Response generation

Best Practices

Data Quality - Clean, structured documents - Validated relationships - Regular updates

Performance - Efficient search algorithms - Caching strategies - Scalable architecture

Evaluation - Accuracy metrics - Relevance assessment - Human validation

Connection to Research

When to Use Each Approach

Plain LLM - General questions - Creative tasks - Quick responses - Limited context needs

RAG - Factual questions - Domain-specific queries - Source citation needs - Up-to-date information

GraphRAG - Complex research questions - Relationship analysis - Historical development - Comprehensive understanding

Research Workflow Integration

Planning Phase - Use GraphRAG for comprehensive literature review - Identify research gaps and relationships - Plan methodology based on existing work

Analysis Phase - Use RAG for specific fact-checking - Apply GraphRAG for relationship analysis - Combine approaches for comprehensive understanding

Writing Phase - Use RAG for source citation - Apply GraphRAG for context development - Ensure comprehensive coverage

Navigation

Previous: [LLM APIs](#) ←

Next: [LLM Showcase](#) →

LLM Showcase and Interface Design

LLM Showcase and Interface Design

Learning Objectives

By the end of this chapter, you will be able to: - Understand principles of good UI design for AI applications - Create user-friendly interfaces for LLM applications - Build complete applications with Gradio - Present and showcase LLM projects effectively - Design interfaces that build trust and usability

Introduction

This final chapter focuses on creating complete, user-friendly applications that showcase your LLM work. You'll learn how to build interfaces that make AI accessible to others and present your research effectively.

Pre-Chapter Learning (90 minutes)

Required Materials:

1. Reading:

- [What is Gradio? \(Gradio Blog\)](#)
Why: Intro to building interfaces for ML models.
- [Principles of Good UI Design \(Nielsen Norman Group\)](#)
Why: Key principles for trustworthy, usable AI interfaces.

2. Video:

- [Build a Simple LLM App with Gradio \(YouTube, 25 mins\)](#)
Why: Quick demo of interface building.
- [GraphRAG \(YouTube, 60 mins\)](#)
Why: Advanced RAG and graph-based retrieval for LLMs.

Session 1: Interface Design Principles (30 minutes)

Why Interface Design Matters for AI

Building Trust - Clear, transparent interfaces - Explainable AI decisions - User control and feedback - Error handling and recovery

Usability - Intuitive navigation - Clear instructions - Appropriate feedback - Accessibility considerations

Research Applications - Making AI tools accessible to participants - Creating research tools for colleagues - Sharing findings with broader audiences - Building collaborative research platforms

Key Design Principles

Visibility of System Status - Show what the AI is doing - Provide progress indicators - Display confidence scores - Explain limitations clearly

User Control and Freedom - Allow users to modify inputs - Provide undo/redo functionality - Enable customization options - Give users choice in outputs

Error Prevention - Validate inputs before processing - Provide helpful error messages - Suggest corrections - Prevent common mistakes

Recognition Rather Than Recall - Use familiar terminology - Provide examples and templates - Show recent interactions - Use consistent design patterns

Session 2: Building with Gradio (45 minutes)

What is Gradio?

Gradio is a Python library that makes it easy to create web interfaces for machine learning models. It's perfect for creating demos and prototypes of AI applications.

Basic Gradio Interface

```
import gradio as gr

def analyze_text(text):
    # Your LLM analysis function
    result = llm_analyze(text)
    return result

# Create interface
iface = gr.Interface(
    fn=analyze_text,
    inputs=gr.Textbox(label="Enter your text"),
    outputs=gr.Textbox(label="Analysis result"),
    title="Text Analysis Tool",
    description="Analyze text using AI"
)

iface.launch()
```

Advanced Gradio Features

Multiple Inputs and Outputs

```
def research_analyzer(text, analysis_type, confidence_threshold):
    # Process multiple inputs
    result = analyze_with_parameters(text, analysis_type, confidence_threshold)
    return result, confidence_score, suggestions

iface = gr.Interface(
    fn=research_analyzer,
    inputs=[
        gr.Textbox(label="Research text"),
        gr.Dropdown(choices=["themes", "sentiment", "topics"], label="Analysis type"),
        gr.Slider(minimum=0, maximum=1, label="Confidence threshold")
    ],
    outputs=[
        gr.Textbox(label="Analysis"),
        gr.Number(label="Confidence"),
        gr.Textbox(label="Suggestions")
    ]
)
```

File Upload and Processing

```
def process_survey_file(file):
    # Process uploaded survey data
    results = analyze_survey_data(file.name)
    return results

iface = gr.Interface(
    fn=process_survey_file,
    inputs=gr.File(label="Upload survey responses"),
    outputs=gr.JSON(label="Analysis results")
)
```

Session 3: Research Application Examples (45 minutes)

Example 1: Survey Analysis Tool

Purpose: Help researchers analyze open-ended survey responses

Features: - Upload CSV files with responses - Choose analysis type (themes, sentiment, topics)
- View results with confidence scores - Export analysis results

Interface Design: - Clear file upload area - Dropdown for analysis options - Progress indicator during processing - Results displayed in organized format

Example 2: Literature Review Assistant

Purpose: Help researchers analyze and synthesize research papers

Features: - Upload PDF papers - Extract key findings and themes - Identify research gaps - Generate synthesis summaries

Interface Design: - Drag-and-drop file upload - Multiple analysis options - Side-by-side comparison view - Export capabilities

Example 3: Interview Analysis Tool

Purpose: Process and analyze interview transcripts

Features: - Upload interview transcripts - Identify key themes and quotes - Generate coding suggestions - Create summary reports

Interface Design: - Text input for transcripts - Real-time analysis - Interactive theme exploration - Export functionality

Session 4: Best Practices for AI Interfaces (30 minutes)

Transparency and Explainability

Show the Process - Display what the AI is doing - Explain how decisions are made - Show confidence levels - Provide alternative interpretations

Clear Limitations - State what the AI cannot do - Explain training data limitations - Show potential biases - Provide disclaimers

User Experience Considerations

Loading States - Show progress indicators - Explain what's happening - Provide estimated completion times - Allow cancellation if needed

Error Handling - Provide clear error messages - Suggest solutions - Offer alternative approaches - Maintain user data

Accessibility - Use clear, readable fonts - Provide keyboard navigation - Include screen reader support - Consider color blindness

Research-Specific Considerations

Data Privacy - Explain data handling - Provide privacy controls - Secure data transmission - Clear data retention policies

Academic Standards - Cite sources and methods - Provide reproducibility information - Include limitations and caveats - Enable peer review

Implementation Guide

Step 1: Define Your Application

Identify Purpose - What research problem does it solve? - Who are the target users? - What are the key features needed? - How will it be used?

Plan the Interface - Sketch the user flow - Identify input and output requirements - Consider user experience - Plan for scalability

Step 2: Build the Backend

Core Functionality - Implement LLM integration - Add data processing - Include error handling - Test thoroughly

API Design - Define clear interfaces - Include proper error handling - Add logging and monitoring - Consider performance

Step 3: Create the Interface

Gradio Setup - Install and configure Gradio - Create basic interface - Add input/output components - Test functionality

User Experience - Add helpful descriptions - Include examples - Provide clear instructions - Test with users

Step 4: Deploy and Share

Deployment Options - Gradio Spaces (free hosting) - Hugging Face Spaces - Local deployment - Cloud platforms

Documentation - Clear usage instructions - API documentation - Research methodology - Limitations and caveats

Connection to Research

When to Build Interfaces

Research Collaboration - Share tools with colleagues - Enable collaborative analysis - Standardize research processes - Facilitate peer review

Public Engagement - Make research accessible - Engage with broader audiences - Demonstrate research impact - Enable public participation

Educational Purposes - Teach research methods - Demonstrate AI capabilities - Provide hands-on learning - Support student research

Best Practices for Research Interfaces

Transparency - Document all methods - Explain AI limitations - Provide source citations - Include disclaimers

Reproducibility - Share code and data - Document parameters - Provide version information - Enable independent verification

Ethics - Consider privacy implications - Address potential biases - Ensure informed consent - Protect participant data

Navigation

Previous: [RAG and Context](#) ←

Next: [Course Conclusion](#) →

References

References

- Abadi, Martín, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, et al. 2016. “Tensorflow: Large-scale machine learning on heterogeneous distributed systems.” *arXiv preprint arXiv:1603.04467*.
- Bender, Emily M, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. “On the dangers of stochastic parrots: Can language models be too big?” *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, 610–623.
- Brown, Tom, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. “Language models are few-shot learners.” *Advances in neural information processing systems* 33: 1877–1901.
- Chowdhery, Aakanksha, and others. 2022. “Palm: Scaling language modeling with pathways.” *arXiv preprint arXiv:2204.02311*.
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding.” *arXiv preprint arXiv:1810.04805*.
- Lewis, Mike, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, and Luke Zettlemoyer. 2020. “Retrieval-augmented generation for knowledge-intensive NLP tasks.” *Advances in Neural Information Processing Systems* 33: 9459–9474.
- Liu, Zhiheng, Xu Zhao, Jingwei Wang, and Jie Tang. 2023. “A survey on large language model based autonomous agents.” *arXiv preprint arXiv:2308.11432*.
- Touvron, Hugo, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. 2023. “Llama: Open and efficient foundation language models.” *arXiv preprint arXiv:2302.13971*.
- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. “Attention is all you need.” *Advances in neural information processing systems* 30.
- Wang, Hongru, and others. 2023. “A survey of large language model based autonomous agents: Architectures, capabilities, and challenges.” *arXiv preprint arXiv:2308.11432*.
- Zhang, Wayne, and others. 2023. “A survey of large language models.” *arXiv preprint arXiv:2303.18223*.
- Zhang, Zhiheng, and others. 2023. “A comprehensive survey on large language model based autonomous agents.” *arXiv preprint arXiv:2308.11432*.